

MODULE - 5 - IA 3

3a)

1a) Explain K-Nearest Neighbour Algorithm.

Ans: * K-Nearest Neighbour is an INSTANCE-BASED METHOD based on LAZY LEARNING where KNN function is approximated locally and all computations are deferred until classification.

* It assumes all instances correspond to points in the n-dimensional space R^n . The nearest neighbours of an instance are defined in terms of the standard Euclidean distance.

* Let an arbitrary instance x be described by a FEATURE VECTOR $((a_1(x), a_2(x), \dots, a_n(x))$ where $a_r(x)$ denotes the value of the r^{th} attribute of instance x .

* Then the distance between two instances x_i & x_j is given by:

$$d(x_i, x_j) = \sqrt{\sum_{r=1}^n (a_r(x_i) - a_r(x_j))^2}$$

25

* In KNN, the target function may be either:

- DISCRETE - VALUED
- REAL - VALUED.

30

5b) Explain KNN algorithm for approximating a discrete-valued function $[f: R^n \rightarrow V]$ with pseudo code.

Camlin	Page
Date	/ /

① KNN for APPROXIMATION OF A DISCRETE VALUED FN.

$$f: R^n \rightarrow V$$

→ TRAINING ALGORITHM:

- * For each TRAINING EXAMPLE $\langle x_i, f(x_i) \rangle$:

 - add the EXAMPLE to the list training-examples

→ CLASSIFICATION ALGORITHM:

- * Given a query instance x_q to be classified:

 - Let x_1, \dots, x_k denote k instances from training-examples that are nearest to x_q .

$$\rightarrow \text{RETURN: } \hat{f}(x_q) \leftarrow \underset{v \in V}{\operatorname{argmax}} \sum_{i=1}^k \delta(v, f(x_i))$$

where $\delta(a, b) = 1$ [if $a = b$]

$\delta(a, b) = 0$ [otherwise]

for 5b)

- * The value $\hat{f}(x_q)$ returned by this algorithm as its estimate of $f(x_q)$ = the most common value of f among the k training examples nearest to x_q .

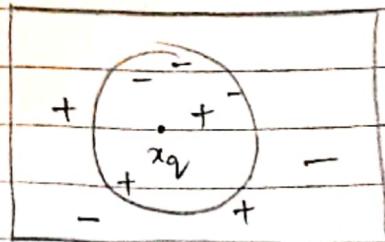
- * If $k=1$, the 1-NN assigns to $\hat{f}(x_q)$ the value of $f(x_i)$

$$\boxed{\hat{f}(x_q) = f(x_i)}$$

 - where x_i is the training instance nearest to x_q

- * For $k = \text{large value}$, the algorithm assigns the most common value among k -nearest training examples

Eg:



+ POSITIVE TRAINING EXAMPLE

- NEGATIVE TRAINING EXAMPLE

1-NN (x_q) → Positive example,

5-NN (x_q) → Negative example.

② KNN for approximation of a REAL-VALUED TARGET FUNCTION $f: \mathbb{R}^n \rightarrow \mathbb{R}$

→ TRAINING ALGORITHM:

* For each TRAINING EXAMPLE $\langle x_i, f(x_i) \rangle$:

⇒ add the example to the list training-examples

→ CLASSIFICATION ALGORITHM:

* Given a query instance x_q to be classified:

⇒ Let x_1, \dots, x_k denote k instances from

training-examples that are nearest to x_q

$$\Rightarrow \text{RETURN: } \hat{f}(x_q) \leftarrow \frac{\sum_{i=1}^k f(x_i)}{k}$$

1b) Describe LOCALLY WEIGHTED REGRESSION (LWR)

Ans: * The term LOCALLY WEIGHTED REGRESSION can be viewed as:

→ LOCAL: Function is approximated based on data near the query point

→ WEIGHTED: Contribution of each training example is weighted by its distance from query point.

→ REGRESSION: Approximation of a REAL-VALUED target function

* Given a new query instance x_q , the general approach is LWR is to construct an approx. \hat{f} that fits the training examples in the neighbourhood, surrounding x_q .

* This approximation is used to calculate the value $\hat{f}(x_q)$ ⇒ estimated target value for query instance

LOCALLY WEIGHTED REGRESSION ALGORITHM.

- * Consider LWR in which TARGET FN, f is approximated near x_q using :

$$\hat{f}(x) = w_0 + w_1 a_1(x) + \dots + w_n a_n(x)$$

where $a_i(x)$ denotes the value of i^{th} attribute of x .

- * Derived methods are used to choose weights that minimize the squared error (using GRADIENT DESCENT)

$$E = \frac{1}{2} \sum_{x \in D} (f(x) - \hat{f}(x))^2$$

$$\Delta w_j = \eta \sum_{x \in D} (f(x) - \hat{f}(x)) a_j(x)$$

constant learning rate

- * Criteria for LOCAL APPROXIMATION :

- ① Minimize the squared error over just the k-NN:

$$E_1(x_q) = \frac{1}{2} \sum_{x \in \text{kNN of } x_q} (f(x) - \hat{f}(x))^2 \quad \dots \dots \quad (1)$$

- ② Minimize the squared error over the entire set D of training examples, while weighting the error of each training example by some decreasing function K of its distance from x_q :

$$E_2() = \frac{1}{2} \sum_{x \in D} (f(x) - \hat{f}(x))^2 K(d(x_q, x)) \quad \dots \dots \quad (2)$$

$(1+2)$

$$E_3(x_q) = \frac{1}{2} \sum_{x \in \text{kNN of } x_q} (f(x) - \hat{f}(x))^2 K(d(x_q, x)) \quad \dots \dots \quad (3)$$

Now,

GRADIENT DESCENT RULE will be :

$$\Delta w_j = \eta \sum_{x \in \text{kNN of } x_q} K(d(x_q, x)) (f(x) - \hat{f}(x)) a_j(x)$$

DISTANCE PENALTY

error summed over only the k -nearest training examples

3a) What is Instance based learning? Explain k-NN algorithm.

Ans: INSTANCE - BASED LEARNING (Memory-based learning)

* It is a family of learning algorithms that, compares new problem instances with instances seen in training, which have been stored in memory, instead of performing explicit generalization.

* Instance-based approaches can construct a different approximation for each distinct query instance that must be classified.

* ~~Time~~ in these algorithms consists of simply storing the presented data. When a new query instance is encountered, a set of similar related instances is retrieved from memory and used to classify a new query instance.

ADVANTAGES

- * Training is very fast
- * Learns complex Target fn.
- * No information loss.

DISADVANTAGES

- * Cost of classification is high
- * Requires large memory
- * Each query initiates a local model formation from scratch.

25
k-NN → Refer 1a)

6a) Explain CADET SYSTEM using Case Based Reasoning.

3b) Explain Case Based Reasoning with example.

Ans: CASE-BASED REASONING (CBR)

* CBR is a learning paradigm based on LAZY LEARNING methods and they classify new query instances by analyzing similar instances while ignoring instances that are very different from the query.

* In CBR, ^{representation of} ~~represent~~ instances are not represented as real valued points, but instead, they use a rich symbolic representation.

* CBR has been applied to problems such as CONCEPTUAL DESIGN of MECHANICAL DEVICES based on a stored library of previous designs, reasoning about NEW LEGAL CASES based on previous rulings and solving PLANNING + SCHEDULING by reusing & combining portions of previous solutions to similar problems.

⇒ EXAMPLE : CADET SYSTEM

* The CADET system employs CBR to assist in the conceptual design of simple mechanical devices such as WATER FAUCETS.

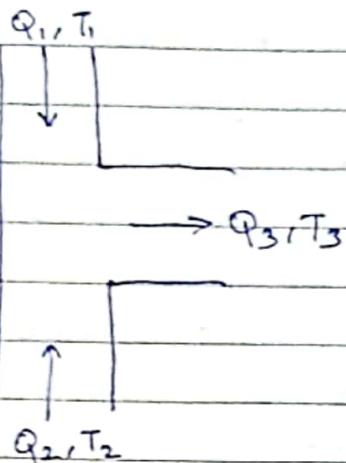
* It uses a library containing approx. 75 previous designs & design fragments to suggest conceptual designs to meet the specifications of new design problems.

* Each instance stored in memory is represented by describing both its structure + qualitative fn.

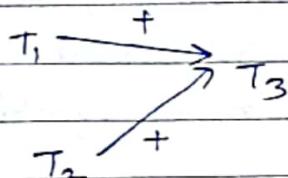
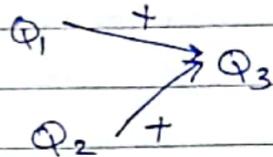
* New design problems are then presented by specifying the desired fn. & requesting the corresponding structure.

I CASE 1: T-JUNCTION PIPE

STORED CASE ($T \rightarrow T_{\text{temp.}}$, $Q \rightarrow \text{waterflow}$)



Function



Notations: $S_1 \xrightarrow{+} S_2$

$\Leftrightarrow S_2$ increases
with S_1

$S_3 \xrightarrow{-} S_4$

S_4 decreases

$T \leftrightarrow Q$ (Qualitative relationship)

Q/H values:

$Q_c \rightarrow$ cold water flow into faucet

$T_c \rightarrow$ temp. of cold water

$Q_h \rightarrow$ hot water flow into faucet

$T_h \rightarrow$ temp. of hot water

$Q_m \rightarrow$ mixed flow out of faucet

$T_m \rightarrow$ temp. of mixed water

CONTROL SIGNALS:

$C_t \rightarrow$ control signal for TEMP. that is input to faucet

CONTROLS

WATER FLOWS

$* C_t \& C_f$ influence Q_c & Q_h
 \Rightarrow indirectly influence

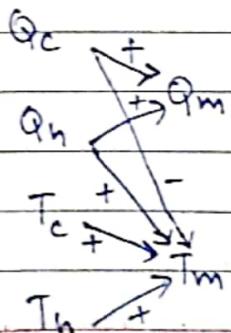
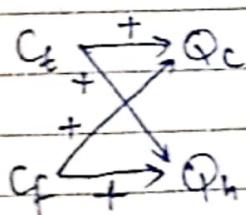
$C_f \rightarrow$ control signal for water flow

output flow $\rightarrow Q_m$

temperature $\rightarrow T_m$

II CASE 2: WATER FAUCET (structure = ?)

Function:



* CADET searches its library for stored cases.

① Match found

\Rightarrow Indicates some STORED case implements DESIRED FN.

\Rightarrow Return case as suggested solution to design problem

② Exact Match not found

CADET may find cases that match various SUBGRAPHS of the DESIRED FN. specification.

4a) Write short notes on RADIAL BASIS FUNCTION (RBF)

(8)

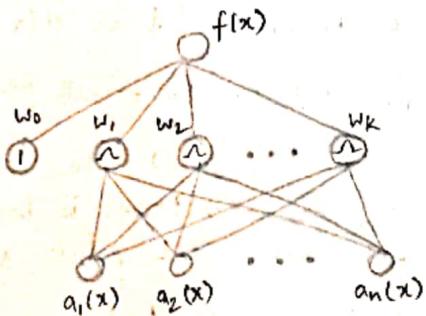
Example ④

RADIAL BASIS FUNCTION (RBF) NETWORK

- * The function given eqn① can be viewed as describing a two layer network where the first layer of units computes the values of the various $K_n(d(x_n, x))$ and where the second layer computes as linear combination of these first-layer unit values.
- * Given a set of training examples of the target fn. RBF networks are typically trained in a two-stage process:
 - ⇒ PROCESS ①: The number k of HIDDEN unit, is determined and each HIDDEN UNIT n is defined by choosing the values of x_n & σ_n^2 , that define its kernel function $K_n(d(x_n, x))$
 - ⇒ PROCESS ②: The weights w , are trained to maximize the fit of the network to the training data, using the global error criterion given by;

$$E \equiv \frac{1}{2} \sum_{x \in D} (f(x) - \hat{f}(x))^2$$

- * Because the kernel fn. are held fixed during this second stage, the linear weight values w , can be trained very efficiently.



* Each hidden unit produces an activation determined by a GAUSSIAN FUNCTION centered at some instance x_n .

* ∴ Its activation will be close to zero unless the input x is near x_n .

* The output unit produces a linear combination of hidden unit activations.

* Multiple output units can also be included ($f_1(x), f_2(x), f_3(x), \dots$)

* supports BACKPROPAGATION

5a) What do you mean by REINFORCEMENT LEARNING? How reinforcement problem differs from other function approximation tasks?

Ans: → REINFORCEMENT LEARNING

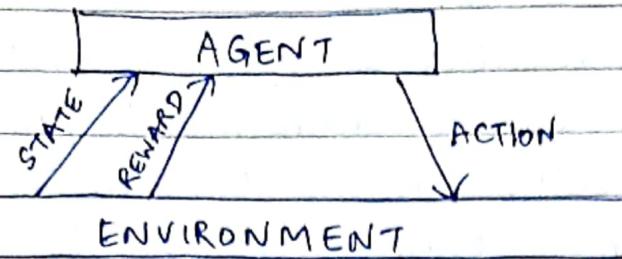
* Reinforcement learning addresses the question of how an AUTONOMOUS AGENT that senses + acts in its ENVIRONMENT can learn to choose OPTIMAL ACTIONS to achieve its GOALS.

Eg: * A learning ROBOT / AGENT observes state of its environment with set of sensors & has set of actions to alter its state.

- * It task is to perform sequences of actions, observe consequences and learn a CONTROL POLICY.
- * CONTROL POLICY chooses actions that maximize the REWARD accumulated over time by the agent.

→ REINFORCEMENT LEARNING PROBLEM

- ① An AGENT interacts with its environment. It exists in an environment described by some set of possible states. (S).
- ② AGENT performs a random set of possible actions A . Each time it performs an action a , in some state s_t , the agent receives a REAL-VALUED REWARD r . (immediate value of STATE-ACTION transition)
- ③ This produces a sequence of states s_i , actions a_i and immediate rewards r_i as shown in figure.
- ④ The agent task is to learn a CONTROL POLICY $T: S \rightarrow A$, that maximizes the expected sum of rewards, discounted exponentially by DELAYS.



$$S_0 \xrightarrow{a_0} S_1 \xrightarrow{a_1} S_2 \xrightarrow{a_2} \dots$$

* CONTROL STRATEGY: $\text{MAX}(r_0 + \gamma r_1 + \gamma^2 r_2 + \dots)$ ($0 \leq \gamma < 1$)
 ↳ discount factor

→ REINFORCEMENT PROBLEM VS FN. APPROXIMATION TASKS

① DELAYED REWARD:

- * TASK: Agent should learn target fn. π that maps
 $S \rightarrow a = \pi(s)$ [but information not available]
 current state optimal action
 in $(s, \pi(s)) \times$
- * provided by only a sequence of immediate reward values
- * Faces TEMPORAL CREDIT ASSIGNMENT problem
 (Determining which is the optimal sequence for eventual / overall rewards)

② EXPLORATION:

- * In RL, the AGENT influences the distribution of training examples by the chosen action sequence.
- * Does EXPERIMENTATION STRATEGY prove to be fruitful?
- * Faces TRADE-OFF in choosing whether to favour EXPLORATION / EXPLOITATION of state-action transitions.

③ PARTIALLY OBSERVABLE STATES:

- * Practically, Agent SENSORS provide only partial information.
- * It needs to reconsider previous observations + current sensor data to choose actions.

④ LIFE - LONG LEARNING:

- * Agent requires to learn several related tasks within same environment, using same sensors.
- * This setting raises the possibility of using PREVIOUSLY OBTAINED EXP./ KNOWLEDGE to reduce sample complexity when learning new tasks.

5. Q-LEARNING

- * Q-learning is a basic form of Reinforcement Learning which uses Q-values (action-values) to iteratively improve the behavior of the learning agent.
- * Q-VALUES / ACTION VALUES: Q-values are defined for states and actions. $Q(S,A)$ is an estimation of how good it is to take the action A at the state S .
- * REWARDS + EPISODES: * An agent over the course of its lifetime starts at a start state, makes a # transitions from its current state to a next state based on its choice of action and also the environment the agent is interacting in.
 - * At every step of transition, the agent from a state takes an action, observes a reward from the environment and then transits to another state.
 - * If at any point of time, the agent ends up in one of the terminating states that means there are no further transitions possible \Rightarrow COMPLETION OF AN EPISODE.
- * TEMPORAL DIFFERENCE / TD-UPDATE: The estimation of $Q(S,A)$ will be iteratively computed using TD-UPDATE RULE. The optimal action in state 's' is the action 'A' that maximizes the sum of the immediate REWARD $R(S,A)$ + the value V^* of the immediate successor state, discounted by γ .

$$\boxed{\pi^*(s) = \underset{A}{\operatorname{argmax}} [R(s,A) + \gamma V^*(\delta(s,A))]} \quad \dots \textcircled{1}$$

- * Q-FUNCTION ANALYSIS: The value of evaluation function $Q(s,a)$ is the reward received immediately upon executing action A from state s with TD-rule:

$$\textcircled{2} \quad [Q(s,A) \equiv R(s,A) + \gamma V^*(\delta(s,A))]$$

\Rightarrow $\textcircled{3}$ from $\textcircled{1}$

$$\boxed{\pi^*(s) \equiv \underset{A}{\operatorname{argmax}} Q(s,A)}$$

* Q-LEARNING ALGORITHM

- The objective of Q-learning is to learn the optimal policy which can inform the agent that what actions should be taken for maximizing the reward under what circumstances. $\pi: S \rightarrow A$
 - The key problem is finding a reliable way to estimate training values for Q, given only a sequence of immediate rewards R spread out over time. This can be achieved by:
 - ⇒ ITERATIVE APPROXIMATION ; $V^*(S) = \max_{A'} Q(S, A')$
 - ⇒ $Q(S, A) = R(S, A) + \gamma \max_{A'} Q(S', A')$ From (2)
- ↳ discounting factor ($0 \leq \gamma < 1$)

* Q-LEARNING ALGORITHM (Pseudo code)

- ① For each S, A initialize the table entry $\hat{Q}(S, A) \leftarrow 0$
- ② Observe the current state S
- ③ Do forever:
 - ① Select an action A and execute it
 - ② Receive immediate reward R
 - ③ Observe new state S'
 - ④ Update the table entry for $\hat{Q}(S, A)$ as:

$$\hat{Q}(S, A) \leftarrow R + \gamma \max_{A'} \hat{Q}(S', A')$$

$$⑤ S \leftarrow S'$$

Eg:

S_1	$R_{66} \rightarrow 73 \rightarrow 100$
	81

$\downarrow A_{RIGHT}$

S_2	$66 \leftarrow R \rightarrow 100$
	81

learner's estimate / hypothesis of actual Q function

- * The agent moves one cell to RIGHT in its grid world ($R=0$) (A_{RIGHT})
- * $\hat{Q}(S, A) \leftarrow R + \gamma \max_{A'} \hat{Q}(S', A')$ [Training rule]
- ⇒ $\hat{Q}(S_1, A_{RIGHT}) \leftarrow R + \gamma \max_{A'} \hat{Q}(S_2, A')$
- ⇒ $\hat{Q}(S_1, A_{RIGHT}) \leftarrow 0 + (0.9) \max \{ 66, 81, 100 \}$
- ⇒ 90 .